

# Parallel Programming

## Exercise Session 9

# Feedback: Exercise 8

# Race Conditions

- Race condition occurs if multiple accesses can happen concurrently and at least one access is a write

Thread 1	Thread 2
x=23;	y=42;
r1 = x;	r2 = y;
v = r1;	w = r2;
z = 2;	

- Thread 1 accesses x, v, z
- Thread 2 accesses y, w
- No race condition, no bad interleaving possible

# Race Conditions

- Race condition occurs if multiple accesses can happen concurrently and at least one access is a write

Thread 1	Thread 2
x = 23;	y = 42;
r1 = x;	r2 = y;
v = r1;	w = r2;
y = 2;	z = 2;

- Thread 1 accesses x, v, **y**
- Thread 2 accesses **y**, w, z
- Both write to y! Result depends on interleaving!

# Relations

- For a set  $S$  we can define a mathematical relation  $R$  for members of  $S$
- Example:
  - Set of natural numbers, relation “greater or equal than”
  - Set of all humans, relation “knows”

# Relations

- Relations can have different properties
- Example:
  - Transitivity:  $a R b$  and  $b R c$  implies  $a R c$
  - The “greater than” relation is transitive.
  - The “knows” relation is not transitive.

# Transitive Closure

- For a relation  $R$ , the smallest relation which contains  $R$  and is transitive, is called the transitive closure of  $R$ .
- Example:
  - For the set of airports in the world, we can define the relation “offers direct flight to”
  - This is (probably) not transitive (show why)
  - The transitive closure also has meaning: “can fly from  $a$  to  $b$  with stops”

# Relations and Code

- When we execute code, “actions” happen, i.e., a variable gets read or written
- We can define relations for these actions, such as “is executed before”
  - Easy to check because it is a local property
  - Can build the transitive closure if we want to know if actions are ordered!
  - Not all actions are ordered!



# Program Order

- Not all actions are ordered!

```
S1: a=23;  
S2: x=3;  
S3: if (x==3) {  
S4:   b += 1;  
      } else {  
S5:   b *= 2;  
S6:   x = 0;  
      }  
S7: x=4;
```

Not in program order!

# Program Order

- Action in mutually exclusive code paths are not in program order
- Actions in different threads are not in program order!
- But ordering was good for proofs!
- Want to allow the compiler / hardware to reorder sometimes for performance.
- Solution: Let compiler reorder whenever it is not “observable” – need to define a subset of special actions which are visible across threads

# Synchronization Actions

- Solution: Let compiler reorder whenever it is not “observable” – need to define a subset of special actions which are visible across threads
  - For this lecture, the most important synchronization actions are
    - Start/End of a thread
    - Read/Write of a volatile or atomic variable
    - Acquire / release of a monitor

# Synchronizes With Relation

- The variable `x` is initially 0.
- Thread A writes `x=5`
- Thread B reads/prints the value of `x`
- We could see 0 -> then we expect that B executed before A
- We could see 5 -> then we expect that A executed before B
- The synchronizes with relation says a read of a volatile must return the last value written to it. It synchronizes with that last write across threads!

# Synchronizes With Relation

- If we combine (the transitive closure of) program order and “synchronizes with” we get the “happens before” order
- Any output/result we see in a Java Program must be consistent with this happens before order

# Java Memory Model Takeaways

- If a variable is not declared volatile you must use the happens-before order to reason about possible values -> requires thought
- If a (primitive) variable is volatile it behaves like an atomic register
- Can sometimes gain performance (and maintain correctness) by not marking everything volatile.
- But you are using Java, is performance really your focus? 😊 - if in doubt declare shared variables as volatile

# Java Memory Model Extra Resources

- The Java Memory Model, by Manson, Pugh and Adve
- Java Language and Virtual Machine Specification

# Lecture Recap



# Atomic operations

- An atomic action is one that effectively happens at once i.e. this action cannot stop in the middle nor be interleaved
- It either happens completely, or it doesn't happen at all.
- No side effects of an atomic action are visible until the action is complete

# Atomic registers

- Atomic registers => support read and write, nothing else
- Usually we think of reads / writes as atomic, i.e., if we write a line such as `x=1` in **pseudocode** we assume it happens atomically and is globally visible.
- This is not true in Java (unless `x` is e.g., `AtomicInteger`)
- `Volatile` makes it globally visible (but not atomic in all cases)

# Atomic registers

- An operation such as  $x++$  (with  $x$  being an atomic register) is NOT atomic!
  - Three steps:  $v = \text{read}(x)$ , increment  $v$ ,  $\text{write}(x, v)$
- Problem with atomic registers:
  - Need  $O(n)$  space to synchronize  $n$  threads  $\rightarrow$  bad
  - Fix: support more than read/write in an atomic operation

# Hardware support for atomic operations

Different atomic operations have been proposed, unclear which is best

- Test-And-Set (TAS)
- Compare-And-Swap (CAS)
- Load Linked / Store Conditional
- <http://docs.oracle.com/javase/tutorial/essential/concurrency/atomic.html>

# Hardware Semantics

**boolean TAS(*memref* s)**

atomic

```
if (mem[s] == 0) {  
    mem[s] = 1;  
    return true;  
}  
else  
    return false;
```

**int CAS (*memref* a, int old, int new)**

atomic

```
oldval = mem[a];  
if (old == oldval)  
    mem[a] = new;  
return oldval;
```

# Locks with atomics

- Now we can implement locks for  $n$  threads using a single variable:
  - Lock: `while (!TAS(l)) {}`
  - Unlock: `mem[l] = 0`

# Bus Contention

- TAS/CAS are read-modify-write operations:
  - Processor assumes we modify the value even if we fail!
  - Need to invalidate cache
  - Threads serialize to read the value while spinning

# TATAS

- Idea: Use normal operation to read first, try TAS only if first read returns 0
- Helps a bit. But what about the case where we see 0 first, then 1 in TAS? Can this happen?
  - Yes, and the more threads the more likely 😊



# Exponential Backoff

- Idea: Each time TAS fails, wait longer until you re-try
- Works well, must tune parameters (how long to wait initially, when to stop increasing)
- Same concept in networks, people talking in a high-latency zoom call, etc.

# Deadlock

- Circular dependency between resources/lock and threads
- Nobody can make progress
- Avoid by introducing global order in which locks are taken
  - Cannot have circles now since all dependencies go “in one direction”
- Or by not using locks at all! (Lock-free, wait-free, more later)

# java.util.concurrent.atomic.AtomicBoolean

```
boolean set();
```

atomically set to value **update** iff current value is **expect**. Return true on success.

```
boolean get();
```

```
boolean compareAndSet(boolean expect, boolean update);
```

```
boolean getAndSet(boolean newValue);
```

sets **newValue** and returns previous value.

# Progress Conditions

- Freedom of deadlock

At least one thread is guaranteed to proceed into the critical section at some point

- Freedom of starvation

All threads are guaranteed to proceed into the critical section at some point

- Freedom of starvation  $\Rightarrow$  Freedom of deadlock

# Peterson Lock

```
let P=1, Q=2; volatile boolean array flag[1..2] = [false, false];  
volatile integer victim = 1
```

## Process P (1)

### loop

**non-critical section**

**flag[P] = true**

**victim = P**

**while(flag[Q] && victim == P);**

**critical section**

**flag[P] = false**

I am  
interested

but you go  
first

We both are  
interested

And you go first

## Process Q (2)

### loop

**non-critical section**

**flag[Q] = true**

**victim = Q**

**while(flag[P] && victim == Q);**

**critical section**

**flag[Q] = false**

# Peterson Lock

- Prove deadlock freedom and starvation freedom
- Hint: Only prove starvation freedom. Deadlock freedom will follow from that.

# Fairness

- Intuitive explanation: If thread A calls `lock()` before thread B, then thread A should enter the critical section (CS) before thread B
- More formally: Divide the protocol into a doorway interval D and a waiting interval W
- D has a finite number of steps
- W has an unbounded number of steps

# Fairness

- So called “first-come-first-served”
- Explanation of notation on blackboard
- Is the Peterson lock fair?
- What is the doorway interval  $D$  and the waiting interval  $W$ ?



# Filter Lock

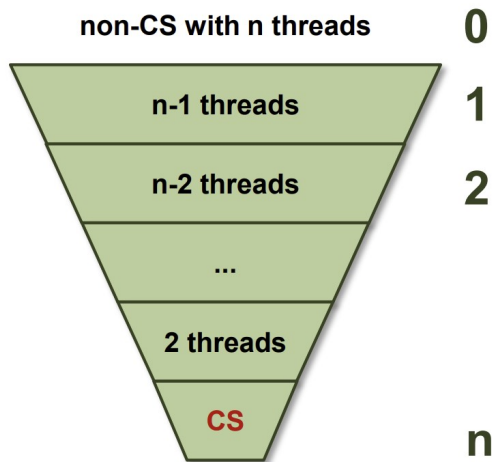
```
int[] level(#threads), int[] victim(#threads)
```

```
lock (me) {  
    for (int i=1; i<n; ++i) {  
        level[me] = i;  
        victim[i] = me;  
        while ( $\exists k \neq me: level[k] \geq i \ \&\& \ victim[i] == me$ ) {};  
    }  
}
```

```
unlock(me) {  
    level[me] = 0;  
}
```

Other threads  
are at same or  
higher level

And I have to wait



# Filter lock is not fair

- What is the doorway section and what is the waiting section?
- Remember the number of steps in the doorway should be bounded
- Doorway section = first two instructions in the first for-loop iteration
- Now find an execution that proves that the filter lock is not fair
- Hint: Use 3 threads

# Assignment 9: Overview

- Analyzing locks
- Atomic operations

# Analyzing locks

- The sample code represents the behavior of a couple that are having dinner together, but they only have a single spoon.
- Prove or disprove that the current implementation provides mutual exclusion.
  - HINT: Use State space diagram

# Atomic operations

- In this task, we will see and analyze:
  - the usage of atomic operations to perform concurrency control, and
  - the cost of using them when having data contention
- For more details, please refer to the assignment sheet

# Kahoot!

<https://create.kahoot.it/details/6c967a18-d238-46fc-9e7b-bcde3d3ac187>